


miniF2F mathd_algebra_109: $(3a + 2b = 12), (a = 4) \vdash b = 0$



Outcome is sparse — it pays once, at $\blacksquare$. The aligned potential is dense: every kernel-verified drop in obligations pays $+\beta$.

So failing early rollouts that get further score higher — within-group variance, hence gradient, where the outcome gives none.

an errored tactic in another rollout pays $-\lambda$